

INFO1105/1905

Data Structures

Week 11: Sorting

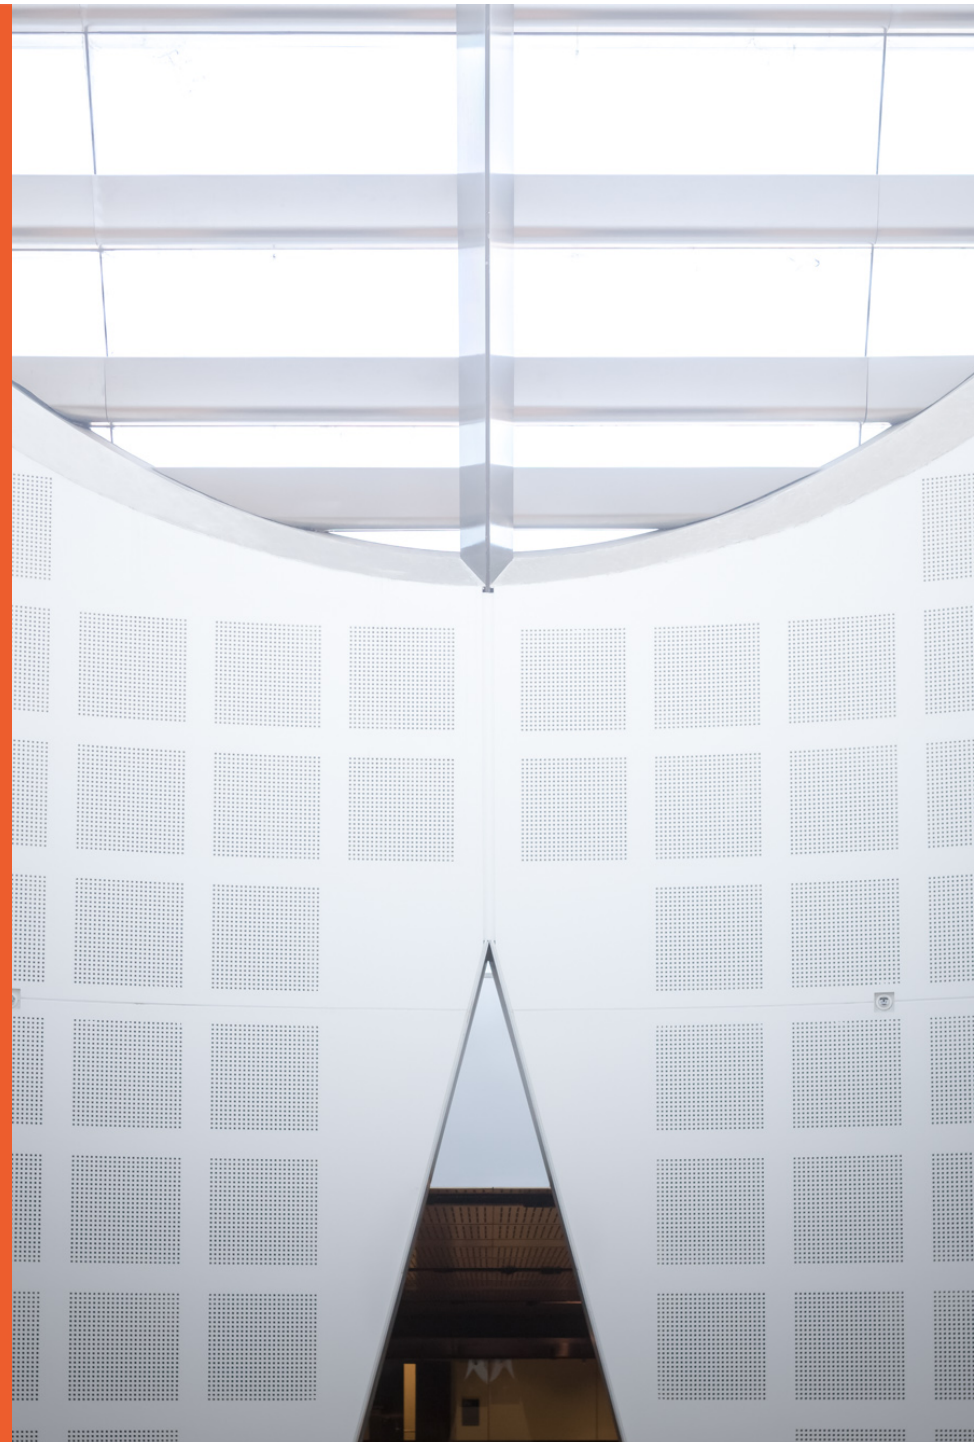
see textbook sections 12.1, 12.2

Professor Alan Fekete

Dr John Stavrakakis

School of Information Technologies

using material from the textbook
and A/Prof Kalina Yacef



Copyright warning

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**).

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.

- These slides contain material from the textbook (Goodrich, Tamassia & Goldwasser)
 - Data structures and algorithms in Java (5th & 6th edition)
- With modifications and additions from the University of Sydney
- The slides are a guide or overview of some big ideas
 - Students are responsible for knowing what is in the referenced sections of the textbook, not just what is in the slides

Reminder! Quiz 5

- Quiz 5 will take place during lab in week 12
- Done online, over a 20 minutes duration,
 - during the last 30 minutes of the lab period, or as indicated by your tutor
- A few multiple choice questions,
 - covering material from lectures of weeks 9, 10, and 11 (labs 10 and 11)
 - hash function and properties
 - separate chaining hashtable
 - open addressing hashtable (linear probing, quadratic probing, double hashing)
 - trie
 - sorting algorithms and their costs

Reminder: Asst 2

- Asst 2 has been released
- Due date postponed to Monday Oct 24 (9pm)
- You must write your own code that implements the interface, *using the data structure described in the instructions*
 - Do not use any Map from other libraries
 - You may use some List types from JCF

Outline

- Sorting algorithms
 - Elementary sorting algorithms based on priority queue (review):
 - insertion sort, selection sort, heapsort
 - Bubblesort
 - Merge-sort
 - Quick-sort
 - Bucket-sort
 - Radix-sort

Recall: Priority Queue Sorting

- We can use a priority queue to sort a set of comparable elements
 1. Insert the elements one by one with a series of **insert** operations
 2. Remove the elements in sorted order with a series of **removeMin** operations
- The running time of this sorting method depends on the priority queue implementation

Algorithm *PQ-Sort*(*S*, *C*)

Input sequence *S*, comparator *C* for the elements of *S*

Output sequence *S* sorted in increasing order according to *C*

P ← priority queue with comparator *C*

while $\neg S.isEmpty()$

e ← *S.removeFirst*()

P.insert(*e*, \emptyset)

while $\neg P.isEmpty()$

e ← *P.removeMin().getKey*()

S.addLast(*e*)

Sequence-based Priority Queue

- Implementation with an unsorted list



- Performance:
 - **insert** takes $O(1)$ time since we can insert the item at the beginning or end of the sequence
 - **removeMin** and **min** take $O(n)$ time since we have to traverse the entire sequence to find the smallest key

- Implementation with a sorted list



- Performance:
 - **insert** takes $O(n)$ time since we have to find the place where to insert the item
 - **removeMin** and **min** take $O(1)$ time, since the smallest key is at the beginning

Selection-Sort

- Selection-sort is the variation of PQ-sort where the priority queue is implemented with an *unsorted* sequence
- Running time of Selection-sort:
 1. Inserting the elements into the priority queue with n **insert** operations takes $O(n)$ time
 2. Removing the elements in sorted order from the priority queue with n **removeMin** operations takes time proportional to
$$1 + 2 + \dots + n$$
- Selection-sort runs in $O(n^2)$ time
- Selection is the bottleneck computation

Selection-Sort Example

	Sequence S	Priority Queue P
Input:	(7,4,8,2,5,3,9)	()
Phase 1		
(a)	(4,8,2,5,3,9)	(7)
(b)	(8,2,5,3,9)	(7,4)
..
(g)	()	(7,4,8,2,5,3,9)
Phase 2		
(a)	(2)	(7,4,8,5,3,9)
(b)	(2,3)	(7,4,8,5,9)
(c)	(2,3,4)	(7,8,5,9)
(d)	(2,3,4,5)	(7,8,9)
(e)	(2,3,4,5,7)	(8,9)
(f)	(2,3,4,5,7,8)	(9)
(g)	(2,3,4,5,7,8,9)	()

Insertion-Sort

- Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted sequence
- Running time of Insertion-sort:
 1. Inserting the elements into the priority queue with n **insert** operations takes time proportional to
$$1 + 2 + \dots + n$$
 2. Removing the elements in sorted order from the priority queue with a series of n **removeMin** operations takes $O(n)$ time
- Insertion-sort runs in $O(n^2)$ time
- Insertion is the bottleneck computation

Insertion-Sort Example

	Sequence S	Priority queue P
Input:	(7,4,8,2,5,3,9)	()

Phase 1

(a)	(4,8,2,5,3,9)	(7)
(b)	(8,2,5,3,9)	(4,7)
(c)	(2,5,3,9)	(4,7,8)
(d)	(5,3,9)	(2,4,7,8)
(e)	(3,9)	(2,4,5,7,8)
(f)	(9)	(2,3,4,5,7,8)
(g)	()	(2,3,4,5,7,8,9)

Phase 2



(a)	(2)	(3,4,5,7,8,9)
(b)	(2,3)	(4,5,7,8,9)
..
(g)	(2,3,4,5,7,8,9)	()

Heap-Sort

- Consider a priority queue with n items implemented by means of a heap
 - the space used is $O(n)$
 - methods `insert` and `removeMin` take $O(\log n)$ time
 - methods `size`, `isEmpty`, and `min` take time $O(1)$ time
- Using a heap-based priority queue, we can sort a sequence of n elements in $O(n \log n)$ time
- The resulting algorithm is called heap-sort
- Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort

Array-based Heap Implementation


- We can represent a heap with n keys by means of an array of length $n + 1$
- For the node at rank i
 - the left child is at rank $2i+1$
 - the right child is at rank $2i + 2$
- Links between nodes are not explicitly stored
- Operation **insert** corresponds to inserting at rank $n + 1$
- Operation **removeMin** corresponds to removing at rank n
- Yields in-place heap-sort



Bottom-up Heap Construction



- Instead of inserting each element one by one into the heap
- We can construct a heap storing n given keys in using a bottom-up construction with $\log n$ phases
- In phase i , pairs of heaps with $2^i - 1$ keys are merged into heaps with $2^{i+1} - 1$ keys



Bubble-sort

- A simple sorting algorithm that is easy to code
- To sort a sequence of n comparable elements
 - Scan the sequence $n-1$ times
 - At each step in a scan, compare the current element with the next and swap them if they are out of order
- Each scan moves the largest remaining element to the end of the sequence
 - the next scan is over a sequence that is one element shorter

Example Bubble-sort

First Pass:

(5 1 4 2 8) \rightarrow (1 5 4 2 8)

(1 5 4 2 8) \rightarrow (1 4 5 2 8)

(1 4 5 2 8) \rightarrow (1 4 2 5 8)

(1 4 2 5 8) \rightarrow (1 4 2 5 8)

Second Pass:

(1 4 2 5 8) \rightarrow (1 4 2 5 8)

(1 4 2 5 8) \rightarrow (1 2 4 5 8)

(1 2 4 5 8) \rightarrow (1 2 4 5 8)

Third Pass:

(1 2 4 5 8) \rightarrow (1 2 4 5 8)

(1 2 4 5 8) \rightarrow (1 2 4 5 8)

Fourth Pass:

(1 2 4 5 8) \rightarrow (1 2 4 5 8)

Bubble-sort algorithm

```
array elements[1..N]
for j:= 1 to N-1 do
    for k := 1 to N-j do
        if elements[k] > elements[k+1] then
            swap(k,k+1, elements)
```

big-Oh Run-time analysis for Bubble-sort

array elements[1..N]

for $j := 1$ **to** $N-1$ **do**

outer loop: n iterations

for $k := 1$ **to** $N-j$ **do**

inner loop: at worst n iterations

if elements[k] > elements[$k+1$] **then**

 swap($k, k+1$, elements)

body of inner loop: constant steps

So, total runtime is at worst $O(n \cdot n \cdot 1) = O(n^2)$

we can do more careful analysis of the inner loop

(which is often a lot less than n iterations

cost is $C \cdot \{(n-1) + (n-2) + \dots + 2 + 1\}$ but this is still $O(n^2)$

Summary of Sorting Algorithms so far

Algorithm	Time	Notes
selection-sort insertion-sort Bubble-sort	$O(n^2)$	<ul style="list-style-type: none">▪ slow▪ in-place▪ for small data sets (< 1K)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ fast▪ in-place▪ for large data sets (1K — 1M)

In-place: uses a small amount of memory in addition to that needed to store the objects being sorted

Divide-and-Conquer

- **Divide-and conquer** is a general algorithm design paradigm:
 - **Divide**: divide the input data S in two disjoint subsets S_1 and S_2
 - **Recur**: solve the subproblems associated with S_1 and S_2
 - **Conquer**: combine the solutions for S_1 and S_2 into a solution for S
- The base case for the recursion are subproblems of size 0 or 1
- **Merge-sort** is a sorting algorithm based on the divide-and-conquer paradigm
- Like heap-sort
 - It has $O(n \log n)$ running time
- Unlike heap-sort
 - It does not use an auxiliary priority queue
 - It accesses data in a sequential manner (suitable to sort data on a disk)

Merge-Sort

- **Merge-sort** on an input sequence S with n elements consists of three steps:
 - **Divide**: partition S into two sequences S_1 and S_2 of about $n/2$ elements each
 - **Recur**: recursively sort S_1 and S_2
 - **Conquer**: merge S_1 and S_2 into a unique sorted sequence

Algorithm *mergeSort*(S)

Input sequence S with n elements

Output sequence S sorted (according to a comparator function)

if $S.size() > 1$

$(S_1, S_2) \leftarrow partition(S, n/2)$

mergeSort(S_1)

mergeSort(S_2)

$S \leftarrow merge(S_1, S_2)$

Merging Two Sorted Sequences

- The conquer step of merge-sort consists of merging two sorted sequences A and B into a sorted sequence S containing the union of the elements of A and B
- Merging two sorted sequences, each with $n/2$ elements and implemented by means of a doubly linked list, takes $O(n)$ time

Algorithm *merge*(A, B)

Input sequences A and B with $n/2$ elements each

Output sorted sequence of $A \cup B$

$S \leftarrow$ empty sequence

while $!A.isEmpty()$ && $!B.isEmpty()$

if $A.first().element() < B.first().element()$

$S.addLast(A.remove(A.first()))$

else

$S.addLast(B.remove(B.first()))$

while $!A.isEmpty()$

$S.addLast(A.remove(A.first()))$


while $!B.isEmpty()$

$S.addLast(B.remove(B.first()))$


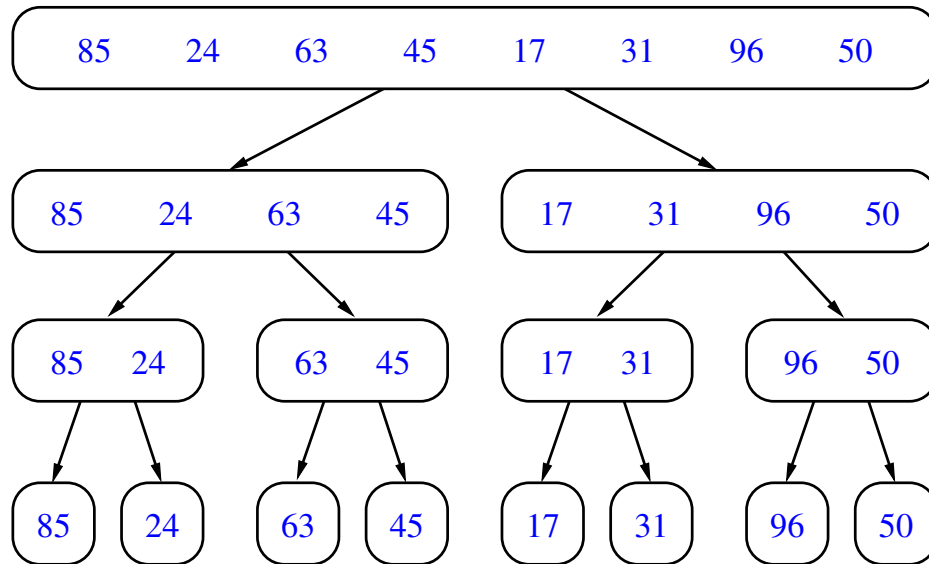
return S

Merge-Sort Tree

- An execution of merge-sort is depicted by a binary tree
 - each node represents a recursive call of merge-sort and stores
 - unsorted sequence before the execution and its partition
 - sorted sequence at the end of the execution
 - the root is the initial call
 - the leaves are calls on subsequences of size 0 or 1




Merge sort trees (input and output sequences)




Execution Example

- Partition



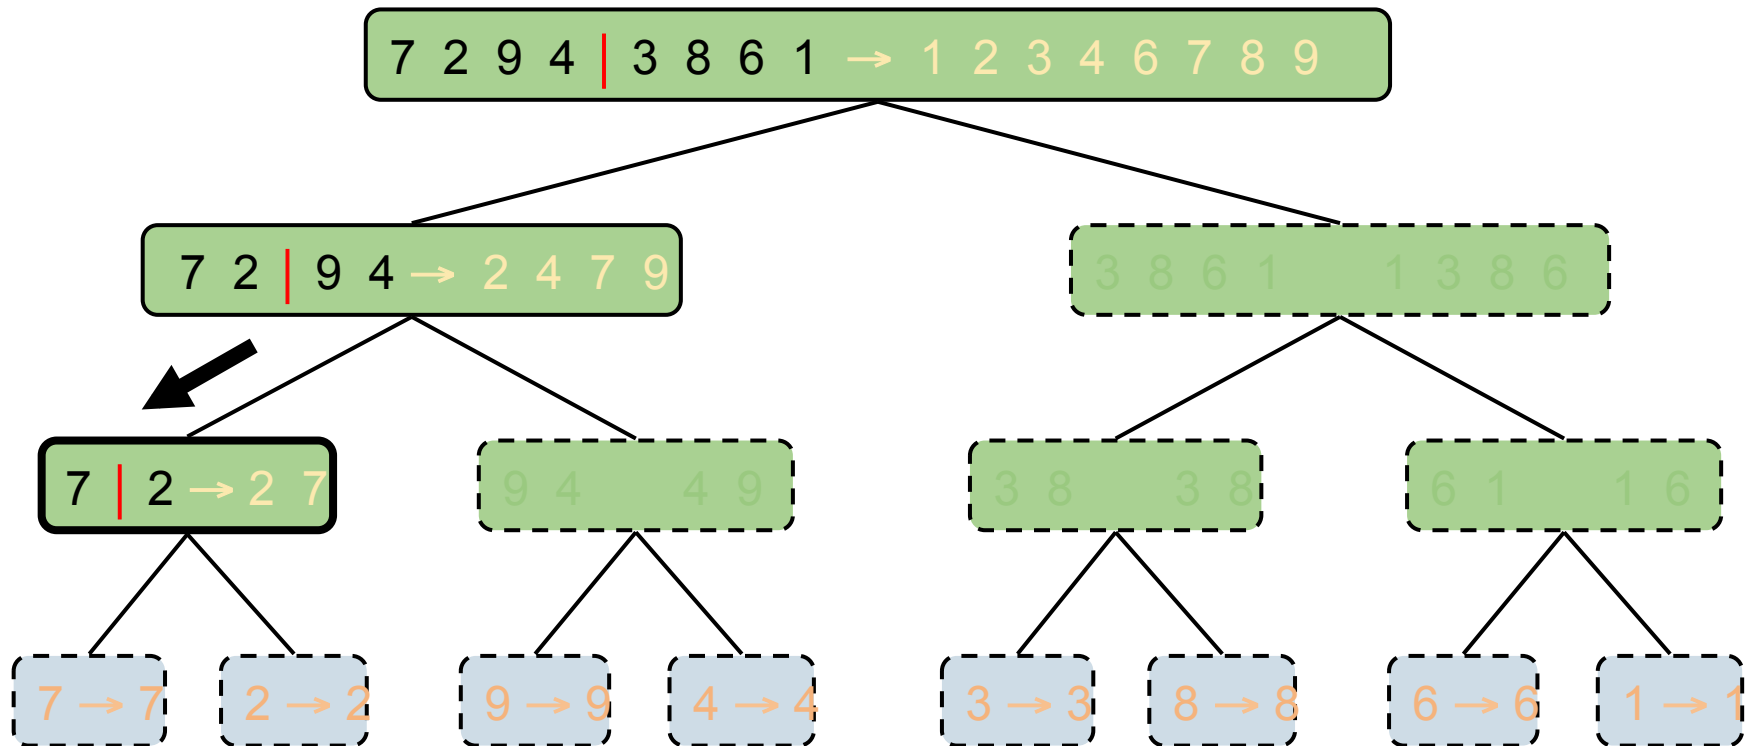
Execution Example (cont.)

- Recursive call, partition



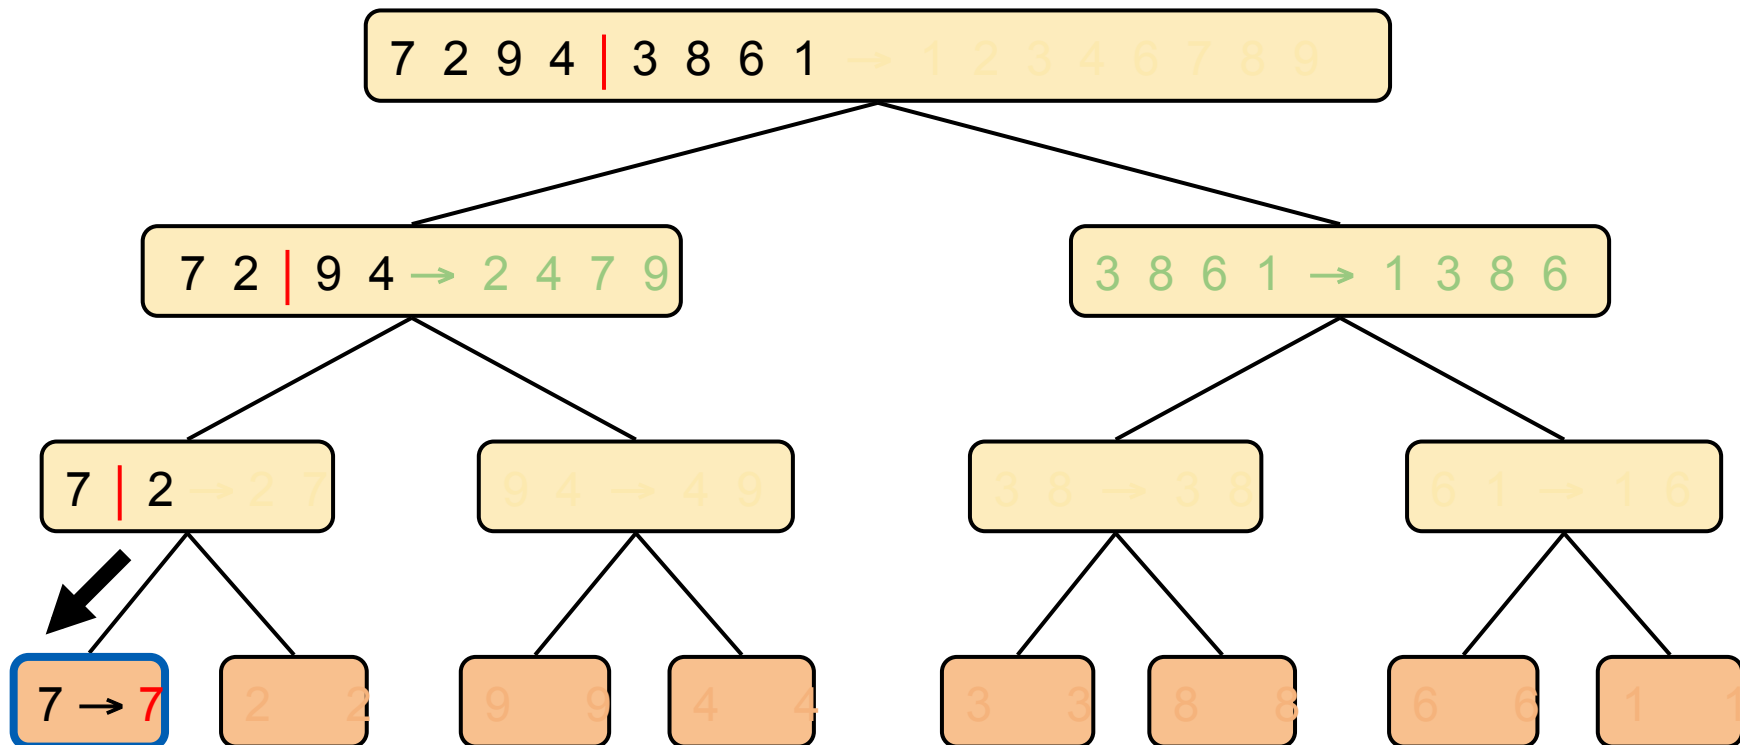
Execution Example (cont.)

- Recursive call, partition



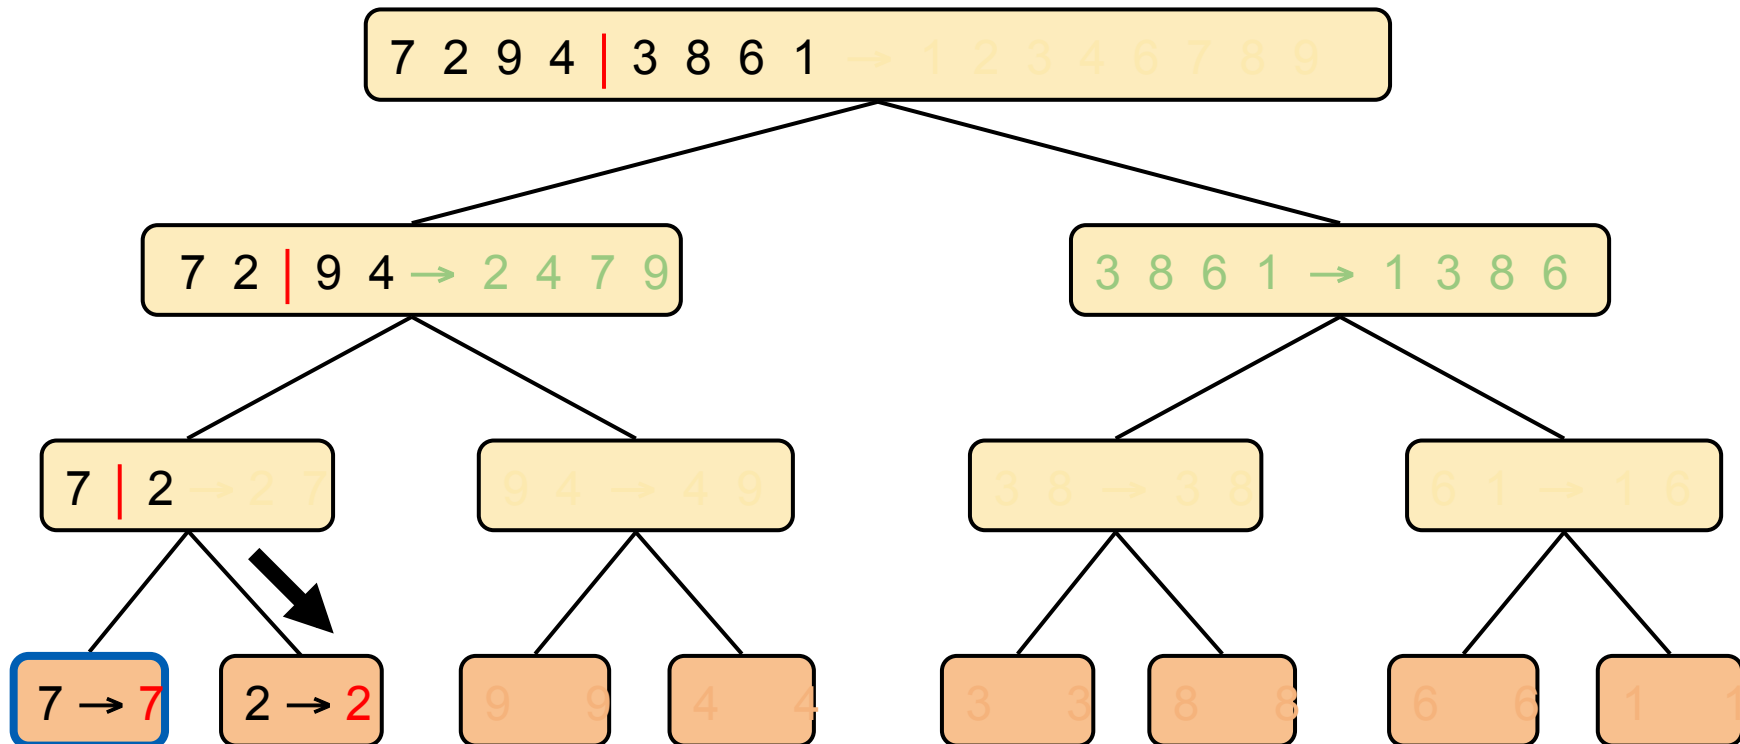
Execution Example (cont.)

- Recursive call, base case



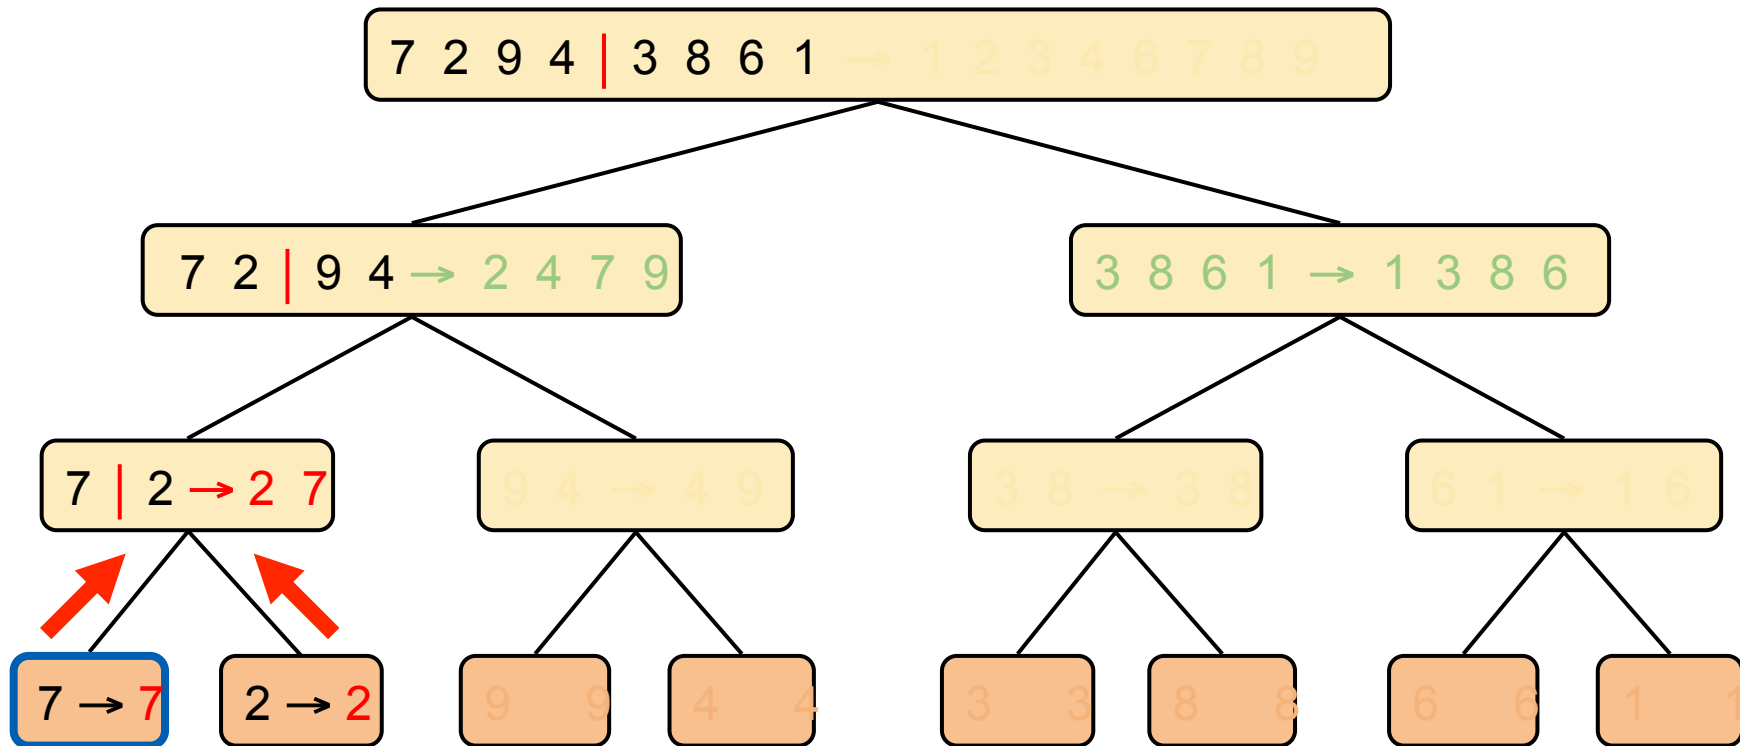
Execution Example (cont.)

- Recursive call, base case



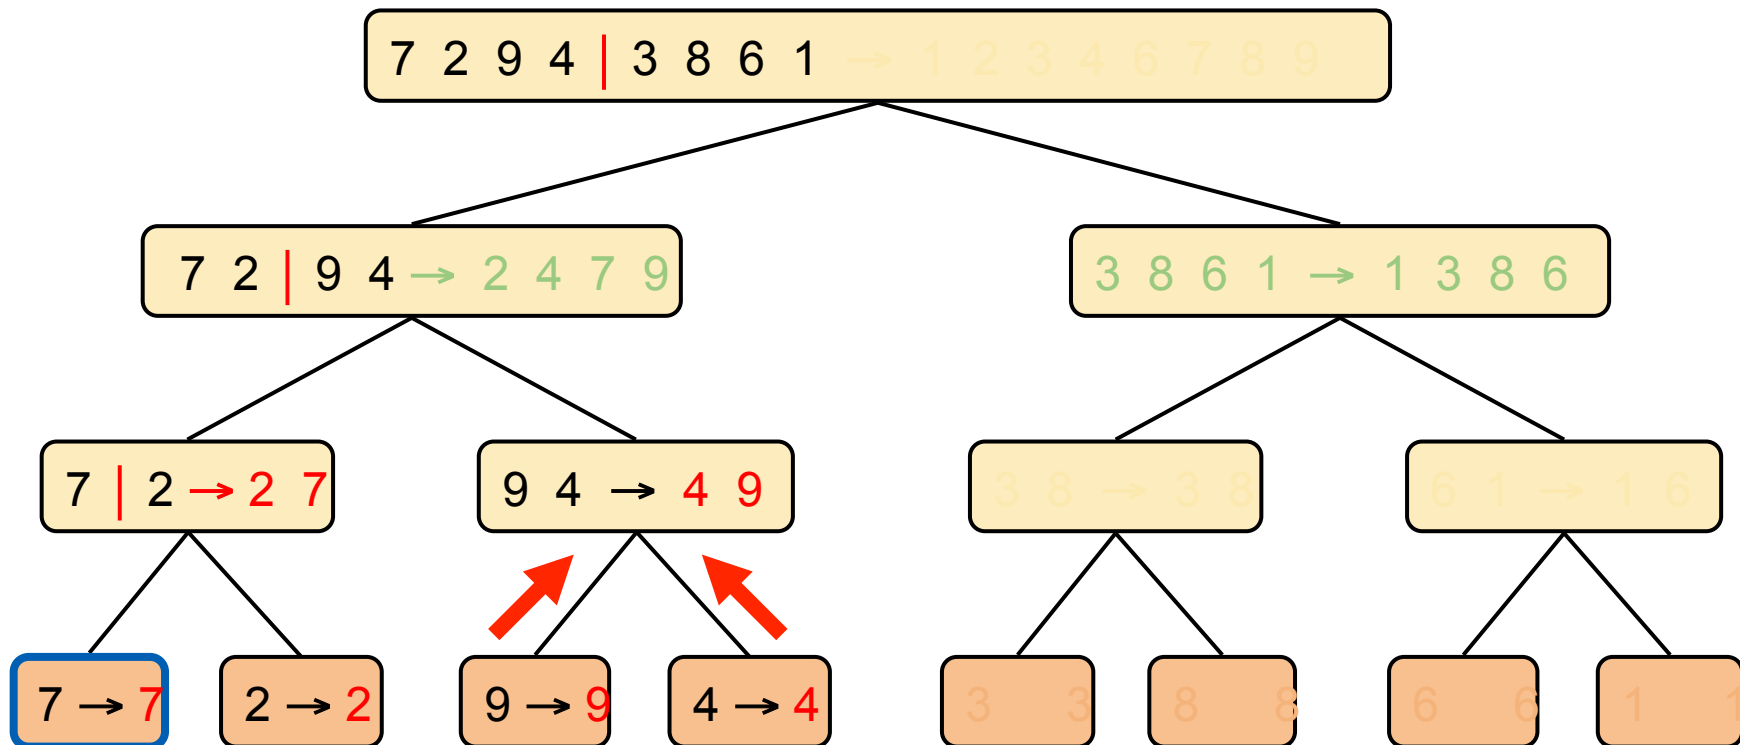
Execution Example (cont.)

- Merge



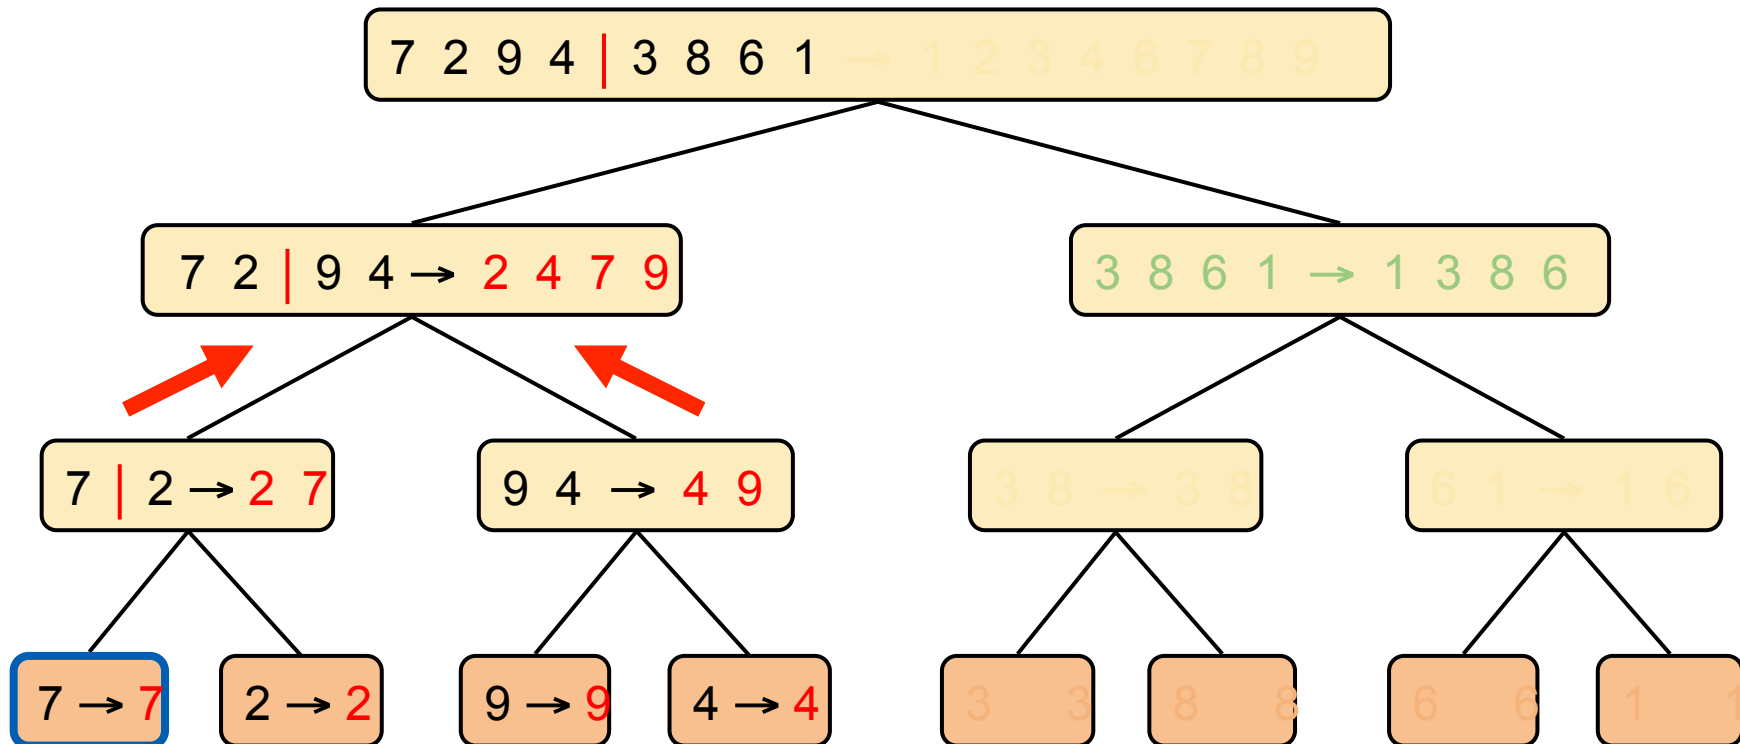
Execution Example (cont.)

- Recursive call, ..., base case, merge



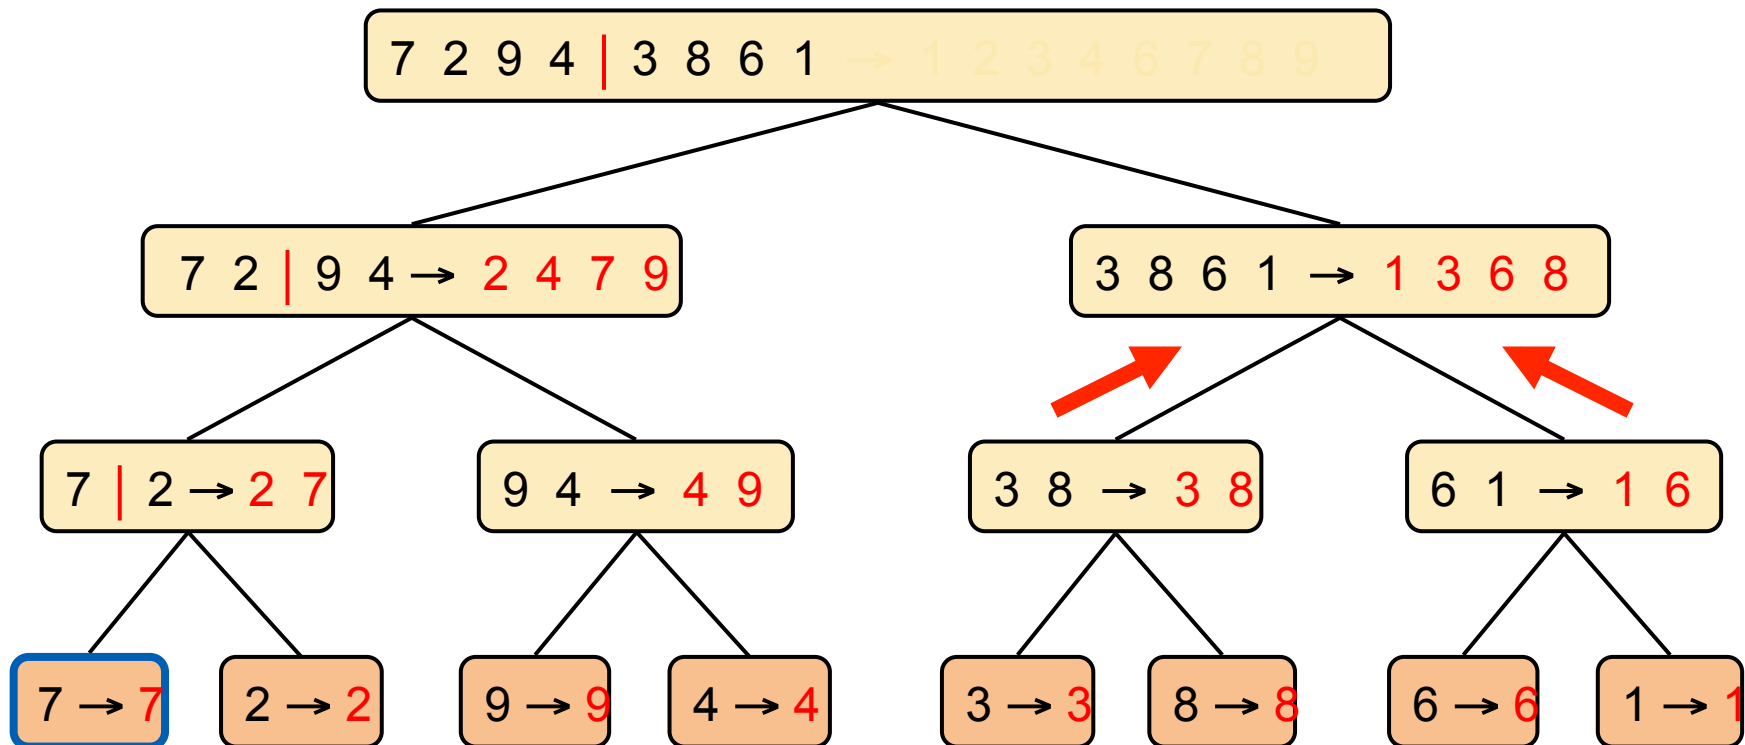
Execution Example (cont.)

- Merge



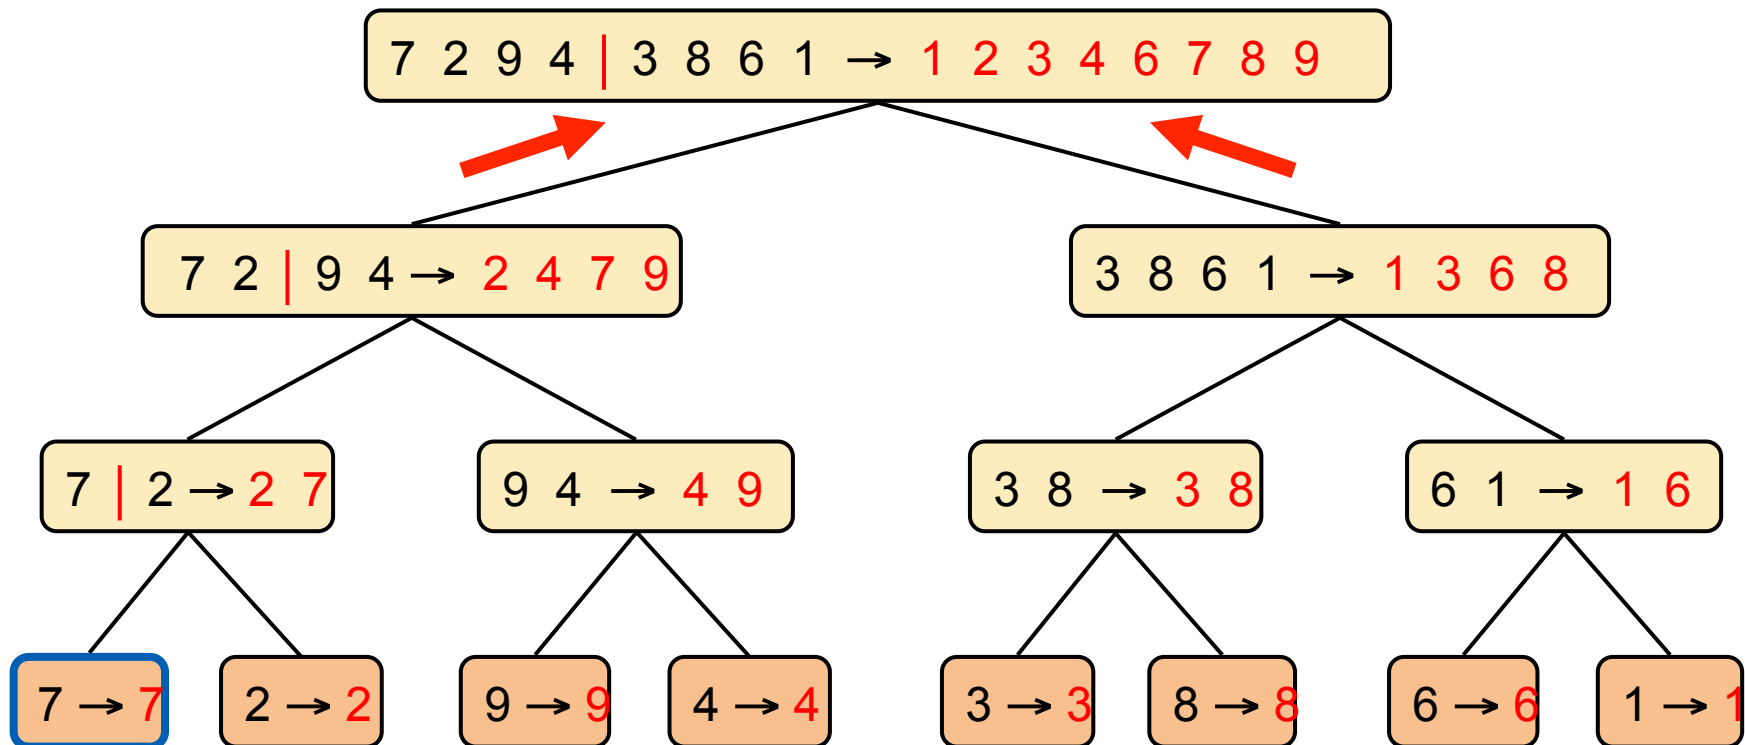
Execution Example (cont.)

- Recursive call, ..., merge, merge



Execution Example (cont.)


- Recursive call, ..., merge, merge



Analysis of Merge-Sort

- The height h of the merge-sort tree is $O(\log n)$
 - at each recursive call we divide in half the sequence,
- The overall amount of work done at the nodes of depth i is $O(n)$
 - we partition and merge 2^i sequences of size $n/2^i$
 - we make 2^{i+1} recursive calls
- Thus, the total running time of merge-sort is $O(n \log n)$

depth	#seqs	size	Time per level
0	1	n	$O(n)$
1	2	$n/2$	$O(n)$
i	2^i	$n/2^i$	$O(n)$
...	$O(n)$



Java Merge Implementation (using arrays)

```

1  /** Merge contents of arrays S1 and S2 into properly sized array S. */
2  public static <K> void merge(K[ ] S1, K[ ] S2, K[ ] S, Comparator<K> comp) {
3      int i = 0, j = 0;
4      while (i + j < S.length) {
5          if (j == S2.length || (i < S1.length && comp.compare(S1[i], S2[j]) < 0))
6              S[i+j] = S1[i++];           // copy ith element of S1 and increment i
7          else
8              S[i+j] = S2[j++];           // copy jth element of S2 and increment j
9      }
10 }

```

	0	1	2	3	4	5	6
S ₁	2	5	8	11	12	14	15

i

	0	1	2	3	4	5	6
S ₂	3	9	10	18	19	22	25

j

	0	1	2	3	4	5	6	7	8	9	10	11	12	13
S	2	3	5	8	9									

i+j

	0	1	2	3	4	5	6
S ₁	2	5	8	11	12	14	15

i

	0	1	2	3	4	5	6
S ₂	3	9	10	18	19	22	25

j





	0	1	2	3	4	5	6	7	8	9	10	11	12	13
S	2	3	5	8	9	10								

i+j






Java Merge-Sort Implementation

```
1  /** Merge-sort contents of array S. */
2  public static <K> void mergeSort(K[ ] S, Comparator<K> comp) {
3      int n = S.length;
4      if (n < 2) return;                // array is trivially sorted
5      // divide
6      int mid = n/2;
7      K[ ] S1 = Arrays.copyOfRange(S, 0, mid);    // copy of first half
8      K[ ] S2 = Arrays.copyOfRange(S, mid, n);    // copy of second half
9      // conquer (with recursion)
10     mergeSort(S1, comp);                // sort copy of first half
11     mergeSort(S2, comp);                // sort copy of second half
12     // merge results
13     merge(S1, S2, S, comp);             // merge sorted halves back into original
14 }
```

Alternative implementation: Linked Lists



Alternative implementation: Linked Lists




Summary of Sorting Algorithms so far

Algorithm	Time	Notes
selection-sort insertion-sort Bubble-sort	$O(n^2)$	<ul style="list-style-type: none">▪ slow▪ in-place▪ for small data sets (< 1K)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ fast▪ in-place▪ for large data sets (1K — 1M)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ fast▪ sequential data access▪ for huge data sets (> 1M)

Advantages of merge sort over heap-sort?


MergeSort parallelises well, stable algorithm




Quick-Sort

Quick-Sort

- **Quick-sort** is a randomized sorting algorithm based on the divide-and-conquer paradigm:
 - **Divide**: pick a random element x (called **pivot**) and partition S into
 - L elements less than x
 - E elements equal x
 - G elements greater than x
 - **Recur**: sort L and G
 - **Conquer**: join L , E and G
- Unlike merge-sort, hard work done *before* the recursive calls



Partition



- We partition an input sequence as follows:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the comparison with the pivot x
- Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$ time
- Thus, the partition step of quick-sort takes $O(n)$ time

Algorithm *partition*(S, p)

Input sequence S , position p of pivot

Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

$L, E, G \leftarrow$ empty sequences

$x \leftarrow S.remove(p)$

while $\neg S.isEmpty()$

$y \leftarrow S.remove(S.first())$

if $y < x$

$L.addLast(y)$

else if $y = x$

$E.addLast(y)$

else $\{ y > x \}$

$G.addLast(y)$

return L, E, G


Java Implementation

```
1  /** Quick-sort contents of a queue. */
2  public static <K> void quickSort(Queue<K> S, Comparator<K> comp) {
3      int n = S.size();
4      if (n < 2) return;           // queue is trivially sorted
5      // divide
6      K pivot = S.first();         // using first as arbitrary pivot
7      Queue<K> L = new LinkedList<>();
8      Queue<K> E = new LinkedList<>();
9      Queue<K> G = new LinkedList<>();
10     while (!S.isEmpty()) {       // divide original into L, E, and G
11         K element = S.dequeue();
12         int c = comp.compare(element, pivot);
13         if (c < 0)                // element is less than pivot
14             L.enqueue(element);
15         else if (c == 0)          // element is equal to pivot
16             E.enqueue(element);
17         else                      // element is greater than pivot
18             G.enqueue(element);
19     }
20     // conquer
21     quickSort(L, comp);           // sort elements less than pivot
22     quickSort(G, comp);           // sort elements greater than pivot
23     // concatenate results
24     while (!L.isEmpty())
25         S.enqueue(L.dequeue());
26     while (!E.isEmpty())
27         S.enqueue(E.dequeue());
28     while (!G.isEmpty())
29         S.enqueue(G.dequeue());
30 }
```

Quick-Sort


Quick-Sort Tree

- An execution of quick-sort is depicted by a binary tree
 - Each node represents a recursive call of quick-sort and stores
 - Unsorted sequence before the execution and its pivot
 - Sorted sequence at the end of the execution
 - The root is the initial call
 - The leaves are calls on subsequences of size 0 or 1




Execution Example

- Pivot selection




Execution Example (cont.)

- Partition, recursive call, pivot selection



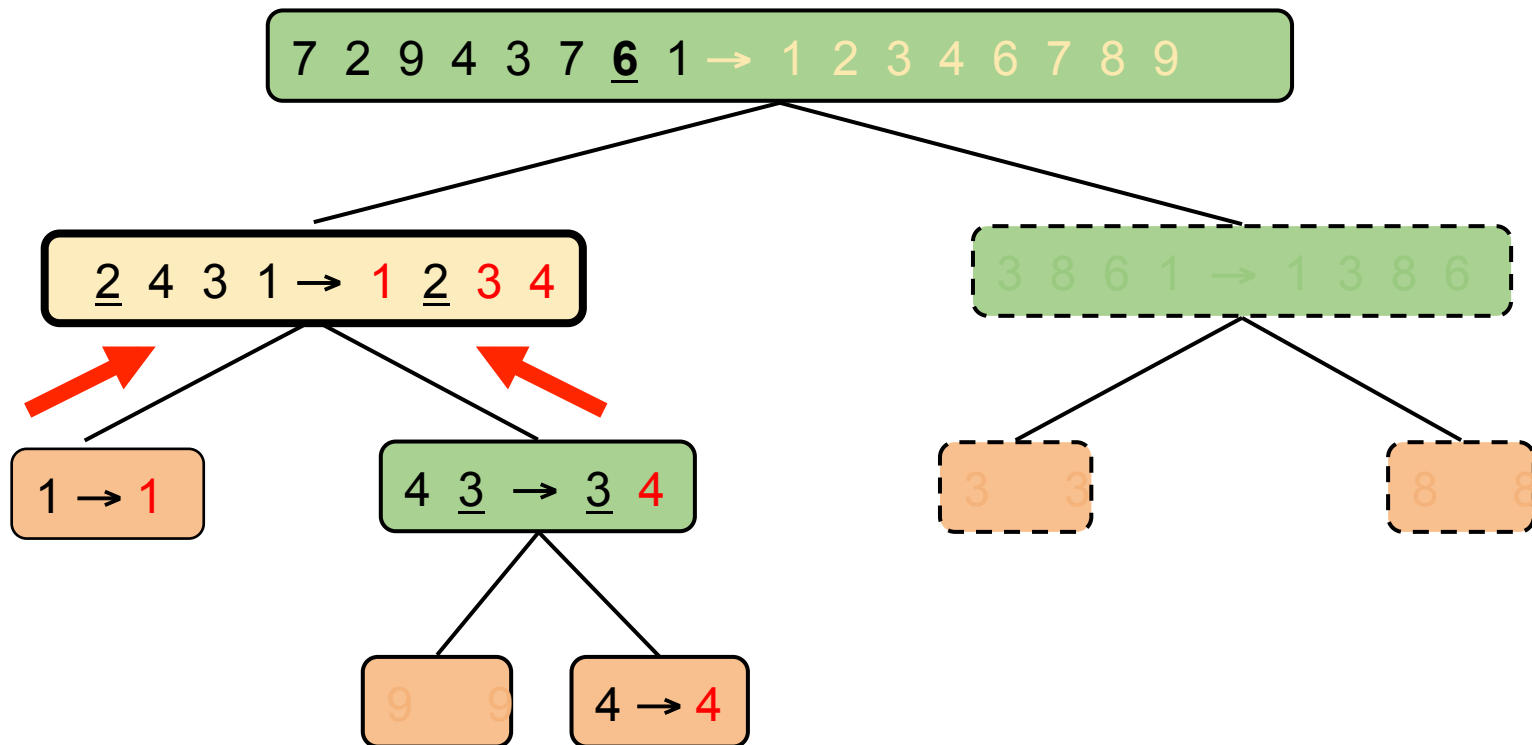
Execution Example (cont.)

- Partition, recursive call, base case



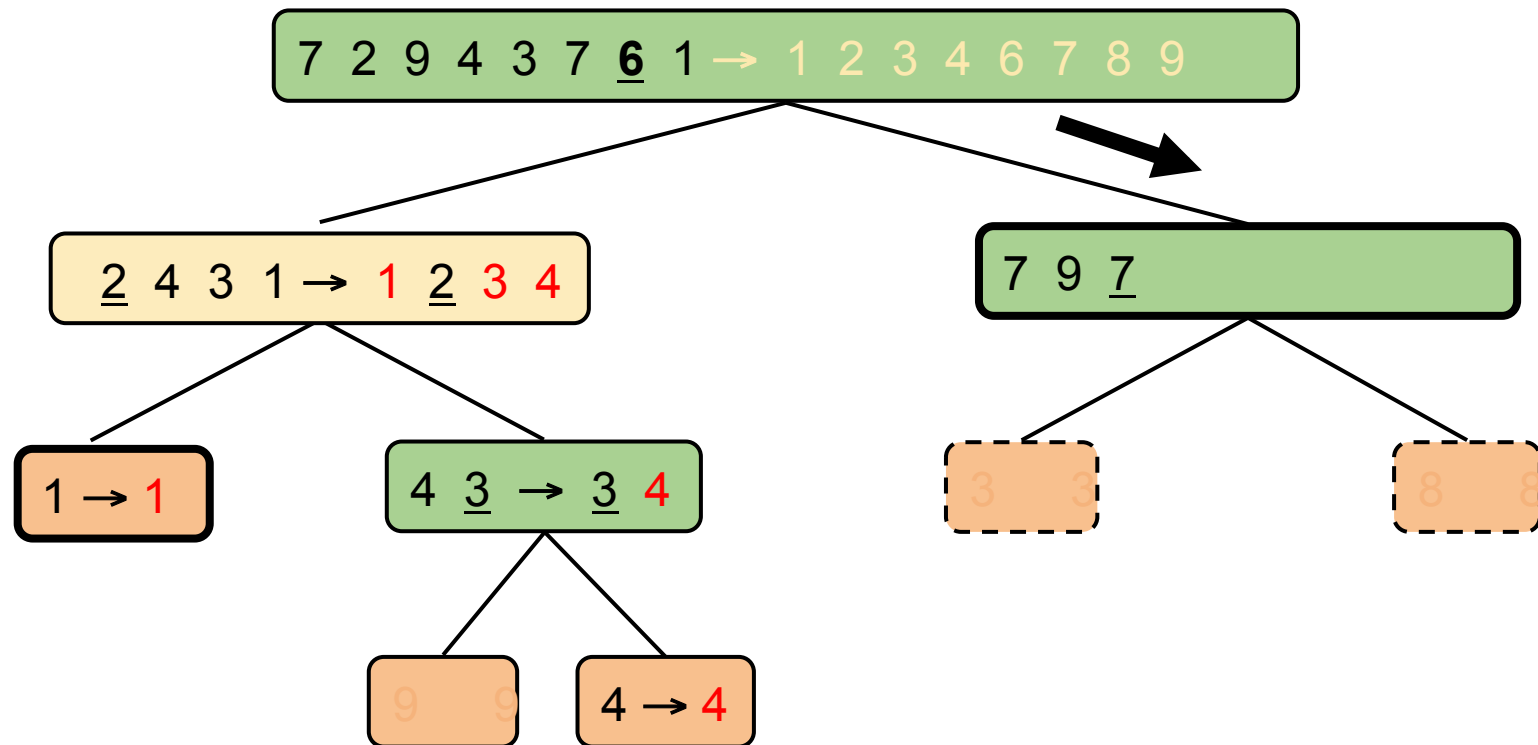
Execution Example (cont.)

- Recursive call, ..., base case, join



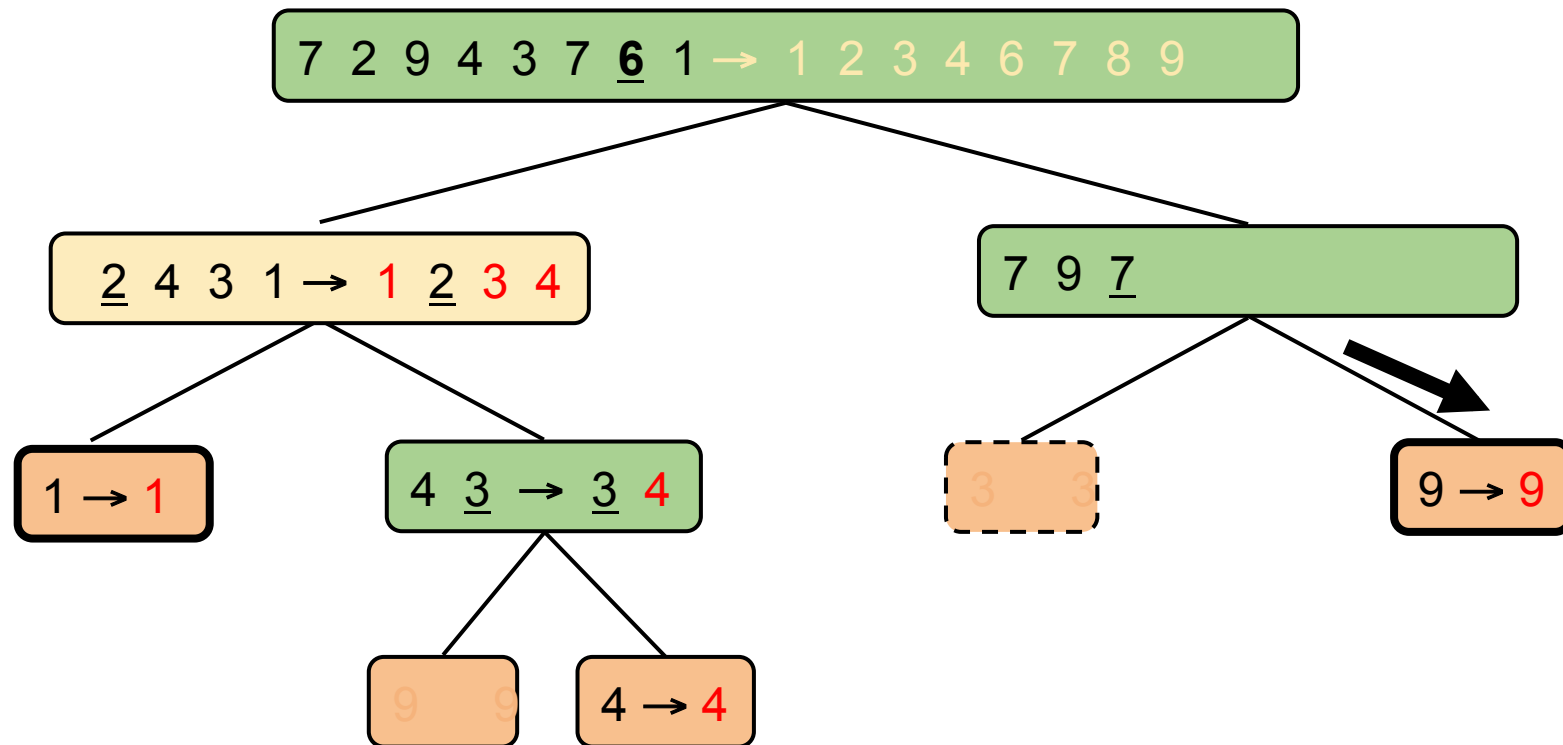
Execution Example (cont.)

- Recursive call, ..., base case, join



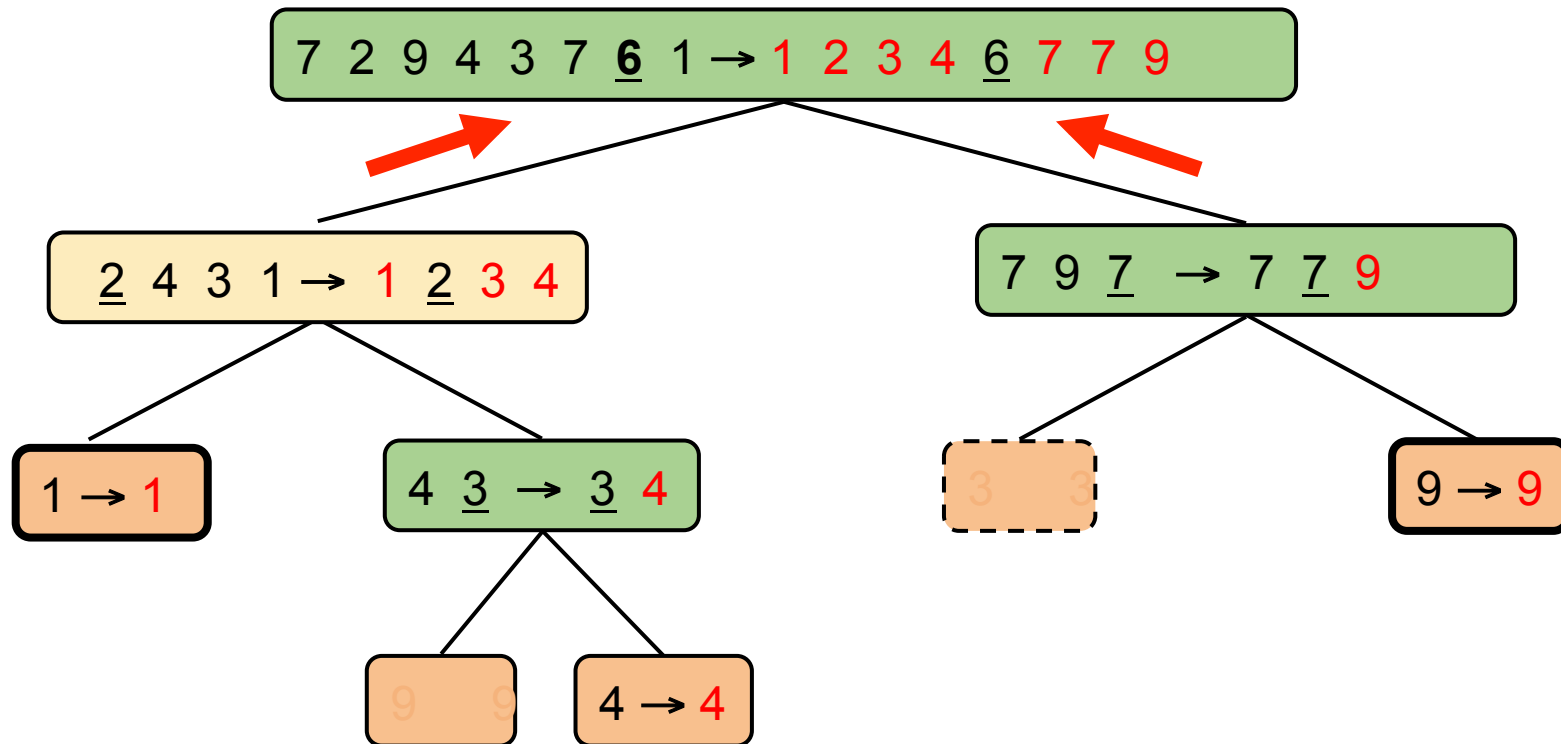
Execution Example (cont.)

- Partition, ..., recursive call, base case




Execution Example (cont.)


- Join, join



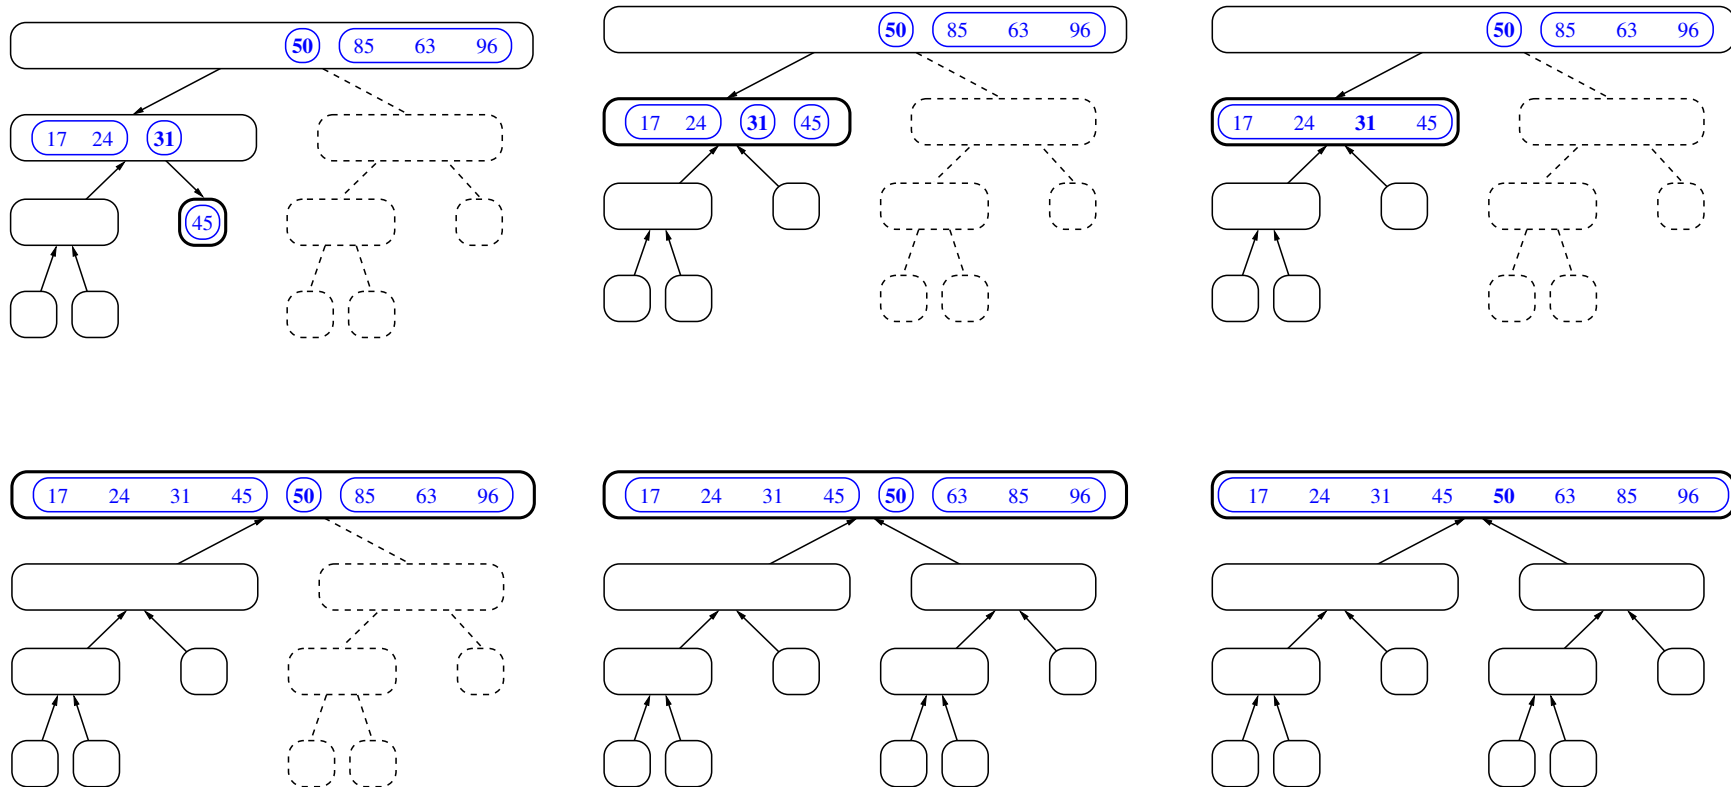
Example 2 (pivot is the last element)



Example 2 (pivot is the last element) cont..




Example 2 (pivot is the last element) cont..



Worst-case Running Time


- The worst case for quick-sort occurs when the pivot is the unique minimum or maximum element
- One of L and G has size $n - 1$ and the other has size 0
- The running time is proportional to the sum
$$n + (n - 1) + \dots + 2 + 1$$
- Thus, the worst-case running time of quick-sort is $O(n^2)$

depth time




Expected Running Time

- Consider a recursive call of quick-sort on a sequence of size s
 - **Good call**: the sizes of L and G are each less than $3s/4$
 - **Bad call**: one of L and G has size greater than $3s/4$



Good call



Bad call


- A call is **good** with probability $1/2$
 - $1/2$ of the possible pivots cause good calls:



Expected Running Time, Part 2

- **Probabilistic Fact:** The expected number of coin tosses required in order to get k heads is $2k$
- For a node of depth i , we expect
 - $i/2$ ancestors are good calls (probability is that every second call is a good call)
 - The size of the input sequence for the current call is at most $(3/4)^{i/2}n$

- Therefore, we have
 - For a node of depth $2\log_{4/3}n$, the expected input size is one
 - The expected height of the quick-sort tree is $O(\log n)$
- The amount of work done at the nodes of the same depth is $O(n)$
- Thus, the **expected** running time of quick-sort is $O(n \log n)$
- **Randomised quick-sort:** picks pivot randomly



In-Place Quick-Sort

- Quick-sort can be implemented to run in-place
- In the partition step, we use replace operations to rearrange the elements of the input sequence such that
 - the elements less than the pivot have rank less than h
 - the elements equal to the pivot have rank between h and k
 - the elements greater than the pivot have rank greater than k
- The recursive calls consider
 - elements with rank less than h
 - elements with rank greater than k



Algorithm *inPlaceQuickSort*(S, l, r)

Input sequence S , ranks l and r

Output sequence S with the elements of rank between l and r rearranged in increasing order

if $l \geq r$

return

$i \leftarrow$ a random integer between l and r

$x \leftarrow S.\text{elemAtRank}(i)$

$(h, k) \leftarrow \text{inPlacePartition}(x)$


inPlaceQuickSort($S, l, h - 1$)

inPlaceQuickSort($S, k + 1, r$)


In-Place Partitioning









- Perform the partition using two indices to split S into L and $E \cup G$ (a similar method can split $E \cup G$ into E and G).




- Repeat until l and r cross:
 - Scan l to the right until finding an element $\geq x$.
 - Scan r to the left until finding an element $< x$.
 - Swap elements at indices l and r



In-Place : divide step



Put pivot in final place



Make recursive calls...


Java Implementation

```
1  /** Sort the subarray S[a..b] inclusive. */
2  private static <K> void quickSortInPlace(K[ ] S, Comparator<K> comp,
3                                          int a, int b) {
4      if (a >= b) return;          // subarray is trivially sorted
5      int left = a;
6      int right = b-1;
7      K pivot = S[b];
8      K temp;                      // temp object used for swapping
9      while (left <= right) {
10         // scan until reaching value equal or larger than pivot (or right marker)
11         while (left <= right && comp.compare(S[left], pivot) < 0) left++;
12         // scan until reaching value equal or smaller than pivot (or left marker)
13         while (left <= right && comp.compare(S[right], pivot) > 0) right--;
14         if (left <= right) {      // indices did not strictly cross
15             // so swap values and shrink range
16             temp = S[left]; S[left] = S[right]; S[right] = temp;
17             left++; right--;
18         }
19     }
20     // put pivot into its final place (currently marked by left index)
21     temp = S[left]; S[left] = S[b]; S[b] = temp;
22     // make recursive calls
23     quickSortInPlace(S, comp, a, left - 1);
24     quickSortInPlace(S, comp, left + 1, b);
25 }
```


Runtimes (experimental)

- Comparisons of mergesort and quicksort (100,000 items)
 - Times (in seconds), List-Based
 - MSort: 55.47
 - QSort: 50.698

mergesort versus quicksort



mergesort versus quicksort



Running times, cpu time

Times (in seconds)

- Array-Based
 - MSort: 0.027 CPU time: 0.0156001
 - QSort: 0.129 CPU time: 0.1248008
- List-Based
 - MSort: 0.732 CPU time: 0.0936006
 - QSort: 0.436 CPU time: 0.0468003

Summary of Sorting Algorithms in this lecture

Algorithm	Time	Notes
selection-sort insertion-sort Bubble-sort	$O(n^2)$	<ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs)
quick-sort	$O(n \log n)$ expected	<ul style="list-style-type: none">▪ in-place, randomized▪ fast (good for large inputs)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ in-place▪ fast (good for large inputs)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ sequential data access▪ fast (good for huge inputs)

Summary

- Read sections 12.1 and 12.2 of textbook (and review section 9.4)
- Sorting algorithms and their costs
 - selection-sort
 - insertion-sort
 - heap-sort
 - bubble-sort
 - merge sort
 - quick sort